

# Practical ML for Engineers

Mathematical Foundations of Supervised Learning

# Learning Objectives

1. Master the notation used in ML
2. Know the different **ingredients** of supervised learning
3. Understand supervised learning as an optimization problem

**ML = Model space + Loss function (+ regularization) + optimization**

## Supervised Learning Setup

We have a (labelled) dataset of input-output examples

$$\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)}) | i = 1, \dots, n\}$$

and we want to find a **model**  $f$  that, given an unseen data pair  $(x, y)$ , performs well at predicting  $y$ :

$$f(x) \approx y$$

## Problems

1. We have **no idea** what  $f$  should look like (**modelization**)
2. We **only have examples** (from the past), but no idea about future  $x$  we might see (**generalization**)

# Model Space and Parameters

- To solve problem nr. 1, and make the model search tractable, we restrict our search to a given **model space**  $\mathcal{M}$ .
- The learning task is then to **find the best**  $f$  in  $\mathcal{M}$ .
- In most cases, our models are parametrized by a finite set of parameters

$$\theta = (\theta_1, \dots, \theta_p) \in \Theta,$$

that belong to a **parameter space**  $\Theta$ .

# Example: Linear Regression Models

## 1D Linear regression model

Model space is the set of all linear functions of  $x$ :

$$\mathcal{M} = \{f(x) = a \cdot x + b \mid a, b \in \mathbb{R}\}.$$

Since the model is completely parametrized by the slope  $a$  and the intercept  $b$ , we can bundle these parameters in a 2d vector, so the parameter space is  $\mathbb{R}^2$ :

$$(a, b) =: \theta \in \Theta := \mathbb{R}^2.$$

For a given parameter  $\theta = (a, b)$ , the corresponding model is

$$f_{\theta}(x) = a \cdot x + b$$

## 1D Polynomial Regression Model

Model space is the set of all polynomial functions of  $x$ , up to degree  $d$ :

$$\mathcal{M} = \{f(x) = a_1 \cdot x + a_2 \cdot x^2 + \dots + a_d \cdot x^d + b \mid a_1, \dots, a_d, b \in \mathbb{R}\}.$$

The model is fully parametrized by the  $(d + 1)$ -dimensional vector:

$$(b, a_1, \dots, a_d) =: \theta \in \Theta := \mathbb{R}^{d+1}.$$



# Model Space: Why bother?

Important

The choice of a **model space**  $\mathcal{M}$  affects subsequent learning in several major ways:

- by restricting the search space, it makes the learning **tractable**
- restricting the search space means that we are also **restricting the class of data that we are allowed to fit** (bias)
- this can, in turn, allow us to incorporate **prior knowledge**
- dimension of the parameter space (complexity) governs the **speed of the learning process** (curse of dimensionality)

## Examples

Model	Functional Form	Parameters	Dimension
Linear	$f_{\theta}(x) = ax + b$	$\theta \in \mathbb{R}^2$	2
Polynomial (degree $p$ )	$f_{\theta}(x) = \sum_{j=0}^p \theta_j x^j$	$\theta \in \mathbb{R}^{p+1}$	$p + 1$
Neural Network	...	Weights $W_i$	...



# Model Space: Influence of complexity

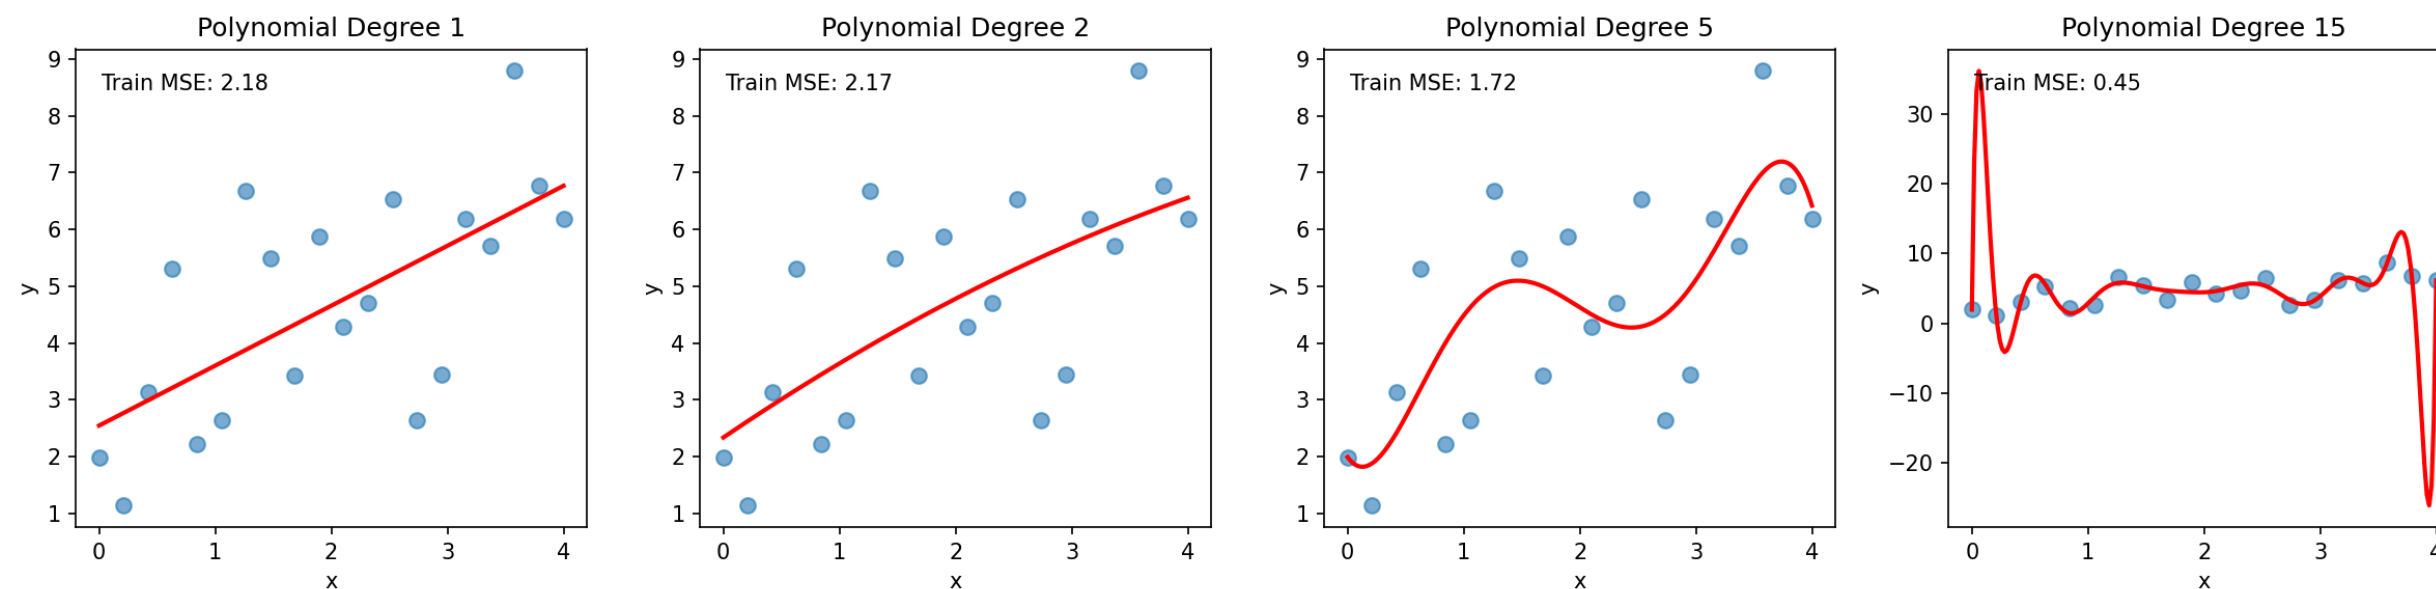
A complex enough model can fit anything.

Example: Polynomial regression

$$f_{\theta}(x) = \theta_1 \cdot x + \theta_2 x^2 + \dots + \theta_n x^n + \theta_{n+1}$$

## 💡 Bias-Variance Trade-off

- **Simple models** (small  $\mathcal{M}$ ): High bias (has difficulty fitting the data), low variance (low sensitivity to noise)
- **Complex models** (large  $\mathcal{M}$ ): Low bias (can fit any data), high variance (too sensitive to data noise)



Will learn later how to **tune model complexity**.

**ML = Model space + Loss function (+ regularization) + optimization**

# Loss Functions

Idea: Quantify how “wrong” our predictions are

- For each example  $(\mathbf{x}, y)$  in the training dataset, we want to know how good our prediction  $\hat{y} = f(\mathbf{x})$  is.
- Then pick the model  $f$  that gives the best predictions.

Features $\mathbf{x}$		Target $y$	$\approx$ ?	Prediction $\hat{y}$
People in Office (Feature 1) $x_1$	Salary (Feature 2) $x_2$	Worked Minutes Week (Target Variable)		Worked Minutes Week (Target Variable)
4	4300 €	2220		2588
12	2700 €	1800		1644
5	3100 €	1920		1870

$\mathcal{D}_{\text{train}}$

- For a parametrized model  $f_{\theta}$ , corresponds to choosing the  $\theta$  that minimizes an “aggregated” loss on the training dataset.

# Loss Functions

- Tells us, for a given prediction  $\hat{y}$  how far we are from the true value  $y$ .
  - Provides a single number to optimize -
- Different cost functions emphasize different aspects

## L2 Loss

$$L(y, \hat{y}) = (y - \hat{y})^2$$

- Penalizes large errors more than small ones
- Influenced by outliers
- Most common loss function for regression

## L1 Loss

$$L(y, \hat{y}) = |y - \hat{y}|$$

- More robust to outliers
- Less statistically well-behaved

# Loss Functions for Learning

Learning = Finding the best point in parameter space

## ! Central Idea: Empirical Risk Minimization

- can sum all losses  $L(y^{(i)}, \hat{y}^{(i)})$  over the training dataset
- for each possible value of the parameter  $\theta$  this give a measure of the “quality” of the model  $f_\theta$

$$J(\theta) = \sum_{i=1}^n L(y^{(i)}, f_\theta(\mathbf{x})^{(i)})$$

- tells us how well the model  $f_\theta$  fits the data for a given  $\theta$
- best model is the one with the smallest loss (smallest risk)


$$\theta^* = \arg \min_{\theta \in \Theta} J(\theta)$$

# Empirical Risk Minimization: Worked Example

## 1D Linear regression

- **Model space:** All linear functions of the form  $a \cdot x + b$
- **Parameters:**  $\theta = (a, b)$
- Each  $(a, b) = (\theta_1, \theta_2)$  pair defines one specific line

Loss Surface in Parameter Space



# Interactive Linear Regression & Loss Surface

## Parameters

$\theta_1$  (Intercept):



0.0

$\theta_2$  (Slope):


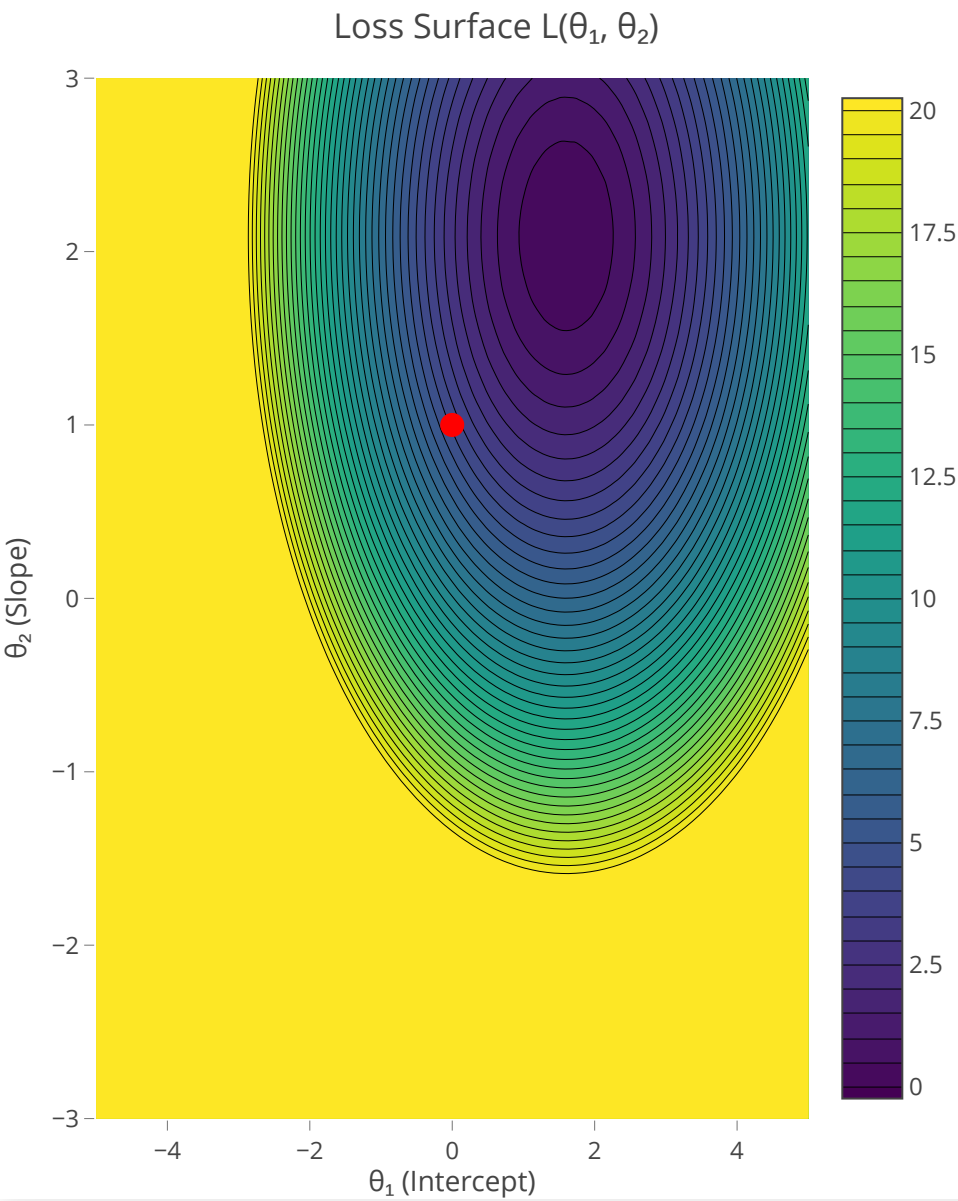


1.0

Model:

$$f(x) = 0.0 + 1.0x$$

Current Loss:  
4.375





## Empirical Risk Minimization: Summary

- (*modelization*) Choose a (parametrized) model class  $f_{\theta} \in \mathcal{M}$
- Choose a loss function  $L(\cdot)$  (L1-loss, L2-loss, ...)
- For each given  $\theta$ , can compute the empirical risk  $\mathcal{J}(\theta)$
- (*model fitting*) Pick the model  $f_{\theta}^*$  that has the minimal risk
- (*prediction*) Use the fitted model  $f_{\theta}^*$  to make predictions at unseen points  $x$

**Question: How can we find the optimal parameter  $\theta^*$  that minimizes the risk?**



## ! Empirical Risk Minimization

- for each possible value of the parameter  $\theta$  define the **empirical risk**:

$$\mathcal{J}(\theta) = \sum_{i=1}^n L(y^{(i)}, f_{\theta}(\mathbf{x})^{(i)})$$

- best model is the one with the smallest loss (smallest risk)

$$\theta^* = \arg \min_{\theta \in \Theta} \mathcal{J}(\theta)$$

How can we find the optimal parameter  $\theta^*$  ?

- Some models have an **analytical solution** (e.g. linear regression)

```
1 ```{python}
2 from sklearn.linear_model import LinearRegression
3 reg = LinearRegression().fit(X, y)
4 y_new = reg.predict(np.array(X_new))
5 ```
```

- In general, perform **iterative minimization** (next chapter). [scipy](#) provides iterative minimization routine:

```
1 ```{python}
2 from scipy.optimize import minimize
3 res = minimize(f, x0,
4               method='nelder-mead',
5               options={'xatol': 1e-8, 'disp': True})
6 ```
```


# How can we choose the Loss function?

# L2 vs L1 Loss Behavior with Outliers

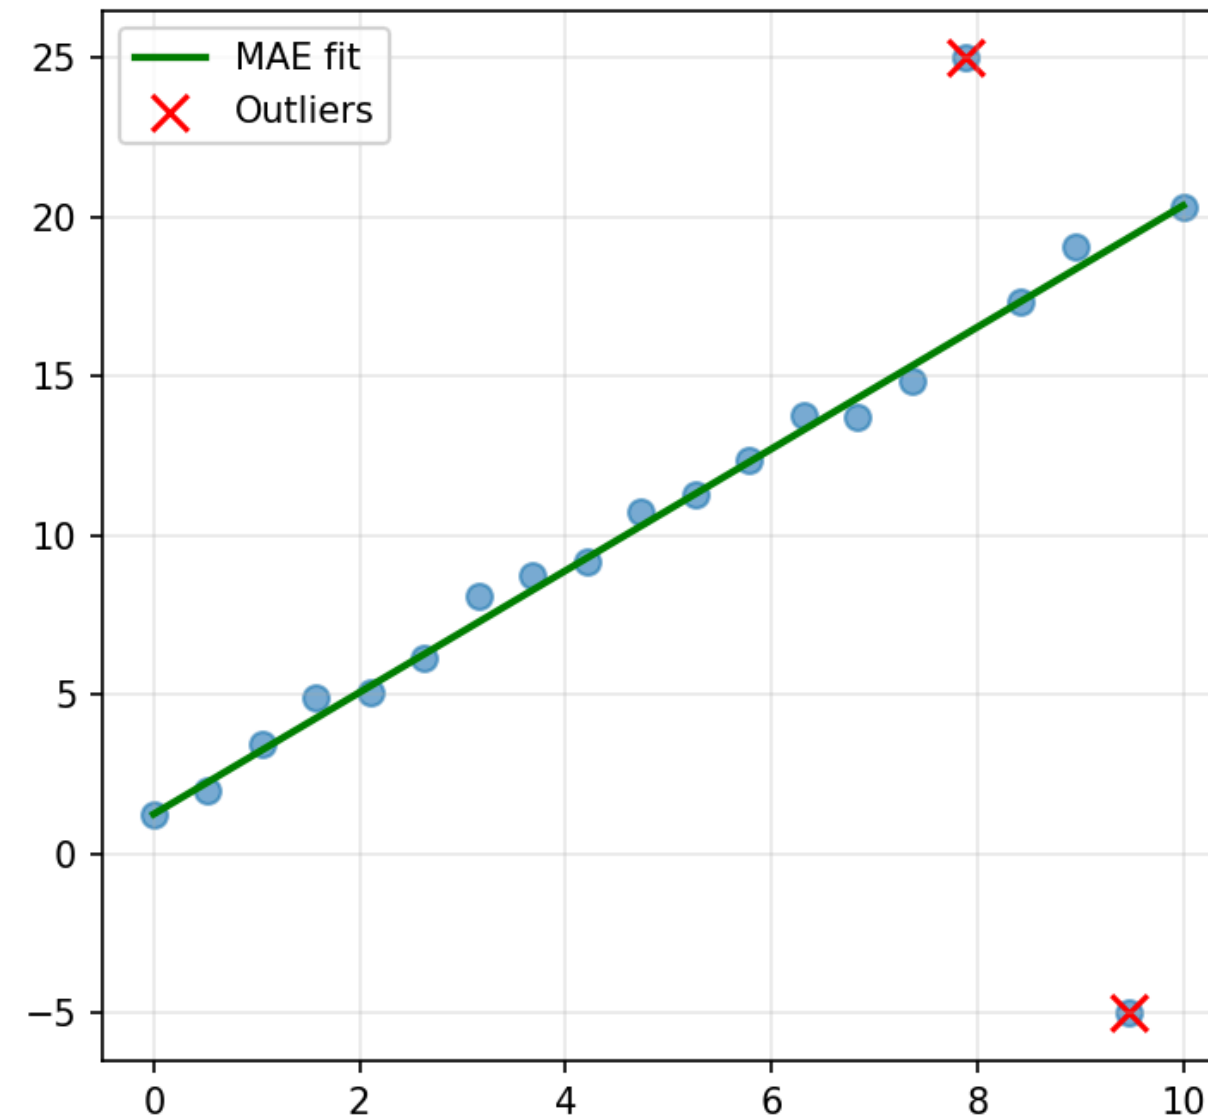
MSE fit:  $y = 1.44x + 2.89$

MAE fit:  $y = 1.91x + 1.25$

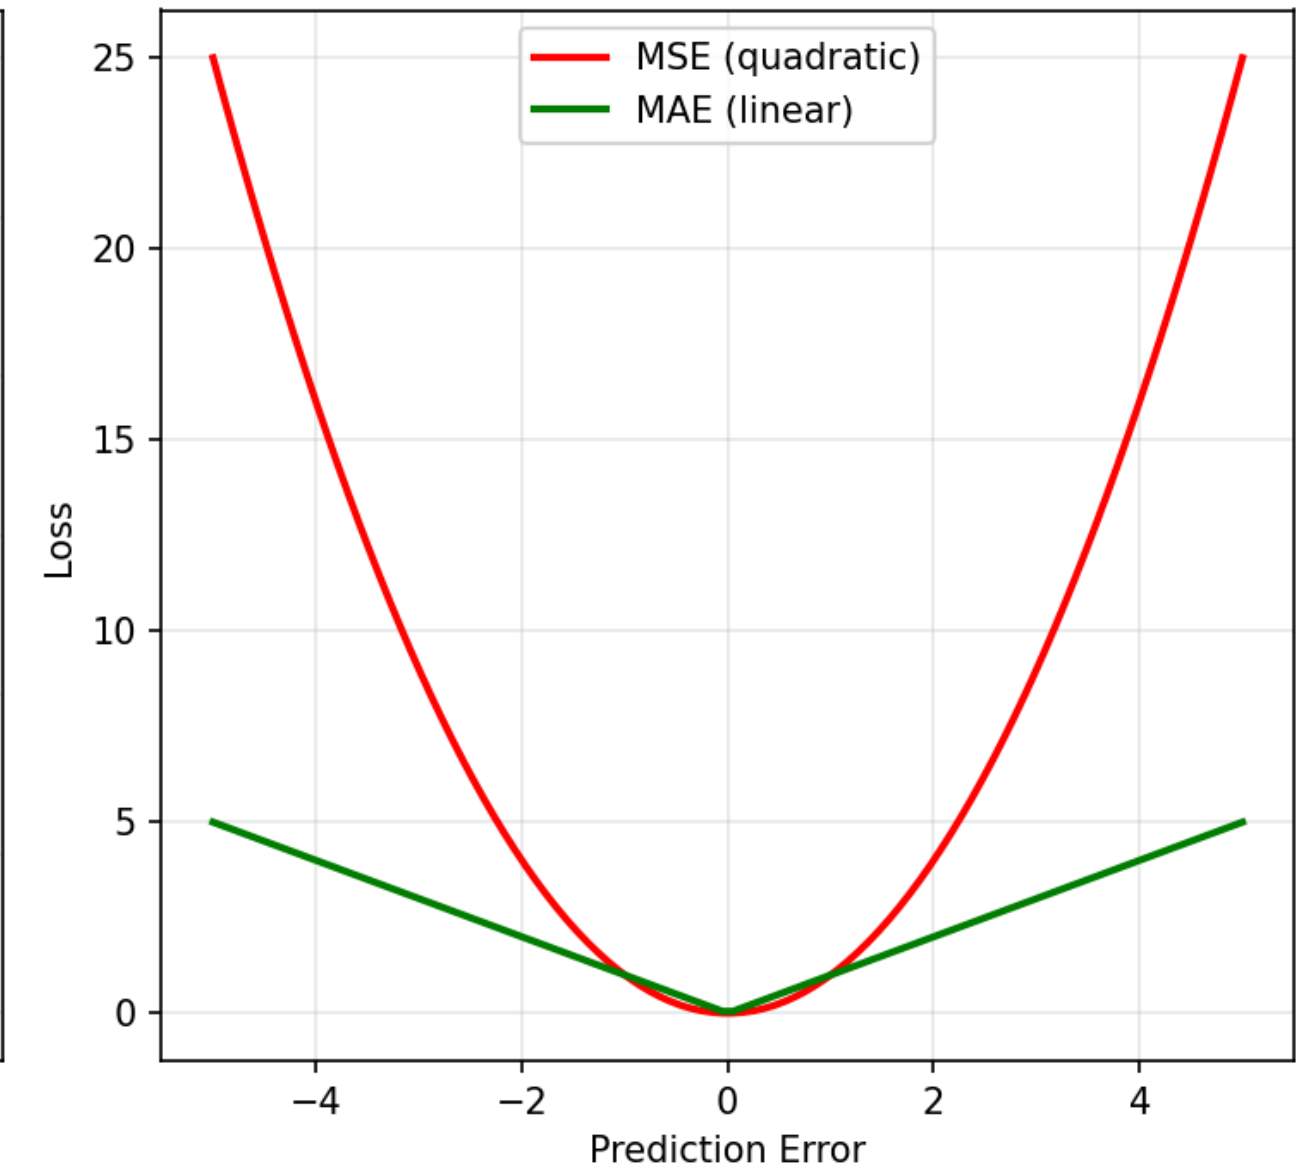
Model with MSE Loss



Model with MAE Loss



Loss Function Comparison



Notice how MSE fit is more influenced by outliers

This is an instance of **overfitting** (more on this in the next chapter).